

Lab report #10
Sadovskaya Veronika

Because Parallel is designed to optimize and speed up big data, a table with 1 million records was created for this lab. Because query results are cached, the table was dropped and recreated for each query.

Task 1 - CREATE Example of Select Parallel execution

1) Select without parallel:

SQL History						
SQL	Connection	TimeStamp	Type	Executed	Duration(seconds)	
SELECT * FROM calendar WHERE day_name IN ('Понедельник','Среда'...	6lab	29.07.22 13:04	SQL	18	0.26	

```

4 EXPLAIN PLAN FOR
5 SELECT * FROM calendar
6 WHERE day_name IN ('Понедельник','Среда', 'Пятница')
7 ORDER BY time_id DESC;

```

Script Output x Query Result x

SQL | All Rows Fetched: 23 in 0,038 seconds

PLAN_TABLE_OUTPUT

8	1	2	1	FX SEND QC (ORDER)		:TQ10001		428K	47M		602	(1)	00:00:01		Q1,01		P->S		QC (ORDER)	
9	3	1	3	SORT ORDER BY				428K	47M	65M	602	(1)	00:00:01		Q1,01		PCWP			
10	4	1	4	FX RECEIVE				428K	47M		599	(1)	00:00:01		Q1,01		PCWP			
11	5	1	5	FX SEND RANGE		:TQ10000		428K	47M		599	(1)	00:00:01		Q1,00		P->P		RANGE	
12	6	1	6	FX BLOCK ITERATOR				428K	47M		599	(1)	00:00:01		Q1,00		PCWC			
13	* 7	1	7	TABLE ACCESS FULL CALENDAR		428K	47M		599	(1)	00:00:01		Q1,00		PCWP					

14

15

16 Predicate Information (identified by operation id):

17 -----

18

19 7 - filter("DAY_NAME"='Понедельник' OR "DAY_NAME"='Пятница' OR "DAY_NAME"='Среда')

20

21 Note

22 -----

23 - Degree of Parallelism is 4 because of table property

2)Select with parallel:

SQL	Connection	TimeStamp	Type	Executed	Duration(seconds)
SELECT /*+ parallel (calendar, 4) */ * FROM calendar WHERE day_nam...	6lab	29.07.22 13:07	SQL	9	0.244

```

10 EXPLAIN PLAN FOR
11 SELECT /*+ parallel (calendar, 4) */ * FROM calendar
12 WHERE day_name IN ('Понедельник', 'Среда', 'Пятница')
13 ORDER BY time_id DESC;

```

Script Output x Query Result x Query Result 1 x Query Result 2 x Query Result 3 x

SQL All Rows Fetched: 23 in 0,047 seconds

PLAN_TABLE_OUTPUT

8	2	FX SEND QC (ORDER)	:TQ10001	428K	47M		602	(1)	00:00:01	Q1,01	P->S	QC (ORDER)
9	3	SORT ORDER BY		428K	47M	65M	602	(1)	00:00:01	Q1,01	PCWP	
10	4	FX RECEIVE		428K	47M		599	(1)	00:00:01	Q1,01	PCWP	
11	5	FX SEND RANGE	:TQ10000	428K	47M		599	(1)	00:00:01	Q1,00	P->P	RANGE
12	6	FX BLOCK ITERATOR		428K	47M		599	(1)	00:00:01	Q1,00	PCWC	
13	* 7	TABLE ACCESS FULL	CALENDAR	428K	47M		599	(1)	00:00:01	Q1,00	PCWP	

16 Predicate Information (identified by operation id):

17 -----

18

19 7 - filter("DAY_NAME"='Понедельник' OR "DAY_NAME"='Пятница' OR "DAY_NAME"='Среда')

20

21 Note

22 -----

23 - Degree of Parallelism is 4 because of table property

As we can see from the results, the query execution time with parallel is less by 0.016 seconds, which almost does not matter. If we look at the execution plans for both queries, we can see that they are identical, that is, when fetching large amounts of data, Oracle itself chose a parallel without being explicitly indicated using hints.

Task 2 - CREATE Example of Parallel DML

1) Without parallel DML

```
19 DELETE FROM calendar
20 WHERE day_name IN ('Понедельник', 'Среда', 'Пятница');
```

Script Output x

Task completed in 1,723 seconds

```

5 -----
6 | 0 | DELETE STATEMENT | | 428K | 7114K | 2153 | (1) | 00:00:01 | | | |
7 | 1 | DELETE | CALENDAR | | | | | | | |
8 | 2 | FX COORDINATOR | | | | | | | |
9 | 3 | FX SEND QC (RANDOM) | :TQ10000 | 428K | 7114K | 2153 | (1) | 00:00:01 | Q1,00 | P->S | QC (RAND)
10 | 4 | FX BLOCK ITERATOR | | 428K | 7114K | 2153 | (1) | 00:00:01 | Q1,00 | PCWC |
11 | 5 | TABLE ACCESS FULL | CALENDAR | 428K | 7114K | 2153 | (1) | 00:00:01 | Q1,00 | PCWF |
12 -----
13
14 Predicate Information (identified by operation id):
15 -----
16
17 5 - filter("DAY_NAME"='Понедельник' OR "DAY_NAME"='Пятница' OR "DAY_NAME"='Среда')
18
19 Note
20 ----
21 - Degree of Parallelism is 4 because of table property
22 - PDML is disabled in current session

```

2) With Parallel DML

```
22 ALTER SESSION ENABLE PARALLEL DML;
23 DELETE /*+ parallel (calendar, 4) */ FROM calendar
24 WHERE day_name IN ('Понедельник', 'Среда', 'Пятница');
25 ALTER SESSION DISABLE PARALLEL DML;
```

Script Output x

Task completed in 0,939 seconds

```

14 PLAN_TABLE_OUTPUT
15 -----
16 | Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time | TQ | IN-OUT | PQ Distrib
17 -----
18 | 0 | DELETE STATEMENT | | 428K | 7114K | 598 (1) | 00:00:01 | | | |
19 | 1 | PX COORDINATOR | | | | | | | | |
20 | 2 | PX SEND QC (RANDOM) | :TQ10000 | 428K | 7114K | 598 (1) | 00:00:01 | Q1,00 | P->S | QC (RAND)
21 | 3 | DELETE | CALENDAR | | | | | | | Q1,00 | PCW
22 | 4 | PX BLOCK ITERATOR | | 428K | 7114K | 598 (1) | 00:00:01 | Q1,00 | PCW
23 | 5 | TABLE ACCESS FULL | CALENDAR | 428K | 7114K | 598 (1) | 00:00:01 | Q1,00 | PCW
24 -----
25
26 Predicate Information (identified by operation id):
27 -----
28
29 5 - filter("DAY_NAME"='Понедельник' OR "DAY_NAME"='Пятница' OR "DAY_NAME"='Среда')
30
31 Note
32 -----
33
34 - Degree of Parallelism is 4 because of table property

```

As we can see, the execution time of a query using parallel is less than half that of a query without parallel. Query plans differ only in the presence of PDML in the current session.

Task 3 - CREATE Example of Parallel DDL

1) Without parallel DDL

```
30 CREATE TABLE calendar AS (SELECT * FROM
31 (SELECT
32 TRUNC( sd + rn ) time_id,
33 TO_CHAR( sd + rn, 'fmDay' ) day_name,
34 TO_CHAR( sd + rn, 'D' ) day_number_in_week,
35 TO_CHAR( sd + rn, 'DD' ) day_number_in_month,
36 TO_CHAR( sd + rn, 'DDD' ) day_number_in_year,
```

Script Output x

Task completed in 23,784 seconds

```

1 PLAN_TABLE_OUTPUT
2
3 Plan hash value: 1220224350
4
5 -----
6 | Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time
7 -----
8 | 0 | CREATE TABLE STATEMENT | | | 1 | 19 | 4 (0) | 00:00:01
9 | 1 | LOAD AS SELECT | CALENDAR | | | | |
10 | 2 | OPTIMIZER STATISTICS GATHERING | | 1 | 19 | 3 (0) | 00:00:01
11 | 3 | VIEW | | | 1 | 19 | 3 (0) | 00:00:01
12 | 4 | COUNT | | | | | |
13 |* 5 | CONNECT BY WITHOUT FILTERING | | | | | |
14 | 6 | FAST DUAL | | 1 | | 3 (0) | 00:00:01
15 -----
16
17 Predicate Information (identified by operation id):
18 -----
19
20 5 - filter(LEVEL<=1000000)

```

2) With parallel DDL (hint)

```
87 CREATE TABLE calendar AS (SELECT /*+ parallel (calendar, 4) */ * FROM
88 (SELECT
89     TRUNC( sd + rn ) time_id,
90     TO_CHAR( sd + rn, 'fmDay' ) day_name,
91     TO_CHAR( sd + rn, 'D' ) day_number_in_week,
92     TO_CHAR( sd + rn, 'DD' ) day_number_in_month
```

PLAN_TABLE_OUTPUT												
9	3	LOAD AS SELECT (HYBRID TSM/HWMB)	CALENDAR							Q1,01	PCWP	
10	4	OPTIMIZER STATISTICS GATHERING			1	19	2	(0)	00:00:01	Q1,01	PCWP	
11	5	PX RECEIVE			1	19	2	(0)	00:00:01	Q1,01	PCWP	
12	6	PX SEND ROUND-ROBIN	:TQ10000		1	19	2	(0)	00:00:01		S->P	RND-ROBIN
13	7	VIEW			1	19	2	(0)	00:00:01			
14	8	COUNT										
15	9	CONNECT BY WITHOUT FILTERING										
16	10	FAST DUAL			1		2	(0)	00:00:01			

19 Predicate Information (identified by operation id):												

22 9 - filter(LEVEL<=1000000)												
23												
24 Note												
25 -----												
26 - Degree of Parallelism is 4 because of table property												

3) With parallel DDL (PARALLEL command)

```
144 CREATE TABLE calendar PARALLEL 4
145 AS (SELECT * FROM
146 (SELECT
147 TRUNC( sd + rn ) time_id,
148 TO_CHAR( sd + rn, 'fmDay' ) day_name,
149 TO_CHAR( sd + rn, 'D' ) day_number_in_week,
150 TO_CHAR( sd + rn, 'DD' ) day_number_in_month,
151 TO_CHAR( sd + rn, 'DDDD' ) day_number_in_year
```

Script Output x

Task completed in 6,655 seconds

PLAN_TABLE_OUTPUT											
9	3	LOAD AS SELECT (HYBRID TSM/HWMB)	CALENDAR						Q1,01	PCWP	
10	4	OPTIMIZER STATISTICS GATHERING			1	19	2	(0)	00:00:01	Q1,01	PCWP
11	5	PK RECEIVE			1	19	2	(0)	00:00:01	Q1,01	PCWP
12	6	PK SEND ROUND-ROBIN	:TQ10000		1	19	2	(0)	00:00:01		S->P RND-ROBIN
13	7	VIEW			1	19	2	(0)	00:00:01		
14	8	COUNT									
15	* 9	CONNECT BY WITHOUT FILTERING									
16	10	FAST DUAL			1		2	(0)	00:00:01		

18											
19 Predicate Information (identified by operation id):											
20 -----											
21											
22 9 - filter(LEVEL<=1000000)											
23											
24 Note											
25 -----											
26 - Degree of Parallelism is 4 because of table property											

As you can see from the query results, the query that runs the fastest is using the PARALLEL command.

The small amount of data, the parallelization does not improve the results. To benefit from parallel query execution, you need to work with a large amount of information.

Task 4 - CREATE Strategy of Parallel execution

The use of parallelism can be a good example of increasing the speed and efficiency of a data warehouse.

Since I have tables in my data warehouse that need to be constantly updated or new data added, such as ORDER_FACT, CLIENT_DIMENSION, EMPLOYEE_DIMENSION, it makes more sense to use the DML parallel to run this process more efficiently at the DW, CL, and SA levels.